

Graph Embeddings

Tim Semenov


timofei@sourced.tech

July 12, 2017

Table of contents

1. Introduction
2. Graph embeddings
3. Experiment results
4. Conclusion

Introduction



Notations

Adjacency matrix:	A
Degree matrix:	$D = \text{diag}(\sum_j A_{ij})$
Transition matrix:	$M = D^{-1}A$
Node neighborhood:	$N(i)$

Notations

Adjacency matrix:	A
Degree matrix:	$D = \text{diag}(\sum_j A_{ij})$
Transition matrix:	$M = D^{-1}A$
Node neighborhood:	$N(i)$
First-order proximity:	edges of adjacency matrix.
Second-order proximity:	similarity between nodes' neighborhoods.

Notations

Pointwise mutual information:

$$PMI(x, y) = \log \frac{P(x, y)}{P(x)P(y)}$$

Positive PMI:

$$PPMI(x, y) = \max(PMI(x, y), 0)$$

Shifted PPMI:

$$SPPMI_k(x, y) = \max(PMI(x, y) - \log k, 0)$$

Graph embedding methods

Taxonomy [2]:

1. Factorization based
2. Random Walk based
3. Deep Learning based
4. Other

Factorization

1. HOPE [4]

- $\min \|S - U^S U^{t^T}\|_F^2$
- $U = [U^S, U^t]$ - the embedding matrix.
- S - proximity matrix:
 - ◊ $S = A^2$ - common neighbors.
 - ◊ $S = ADA$ - Adamic-Adar.
 - ◊ ...

2. GraRep [1]

- $X_k = SPPMI_\beta(A^k)$, $k = 1, \dots, K$
- $[U_k, \Sigma_k, V_k^T] = SVD(X_k)$
- $W_k = U_{k,d}(\Sigma_{k,d})^{\frac{1}{2}}$

Random Walk: node2vec

Optimization problem [3]:

$$\max \sum_{v_i \in V} \log P(N(i)|v_i)$$

Conditional independence:

$$P(N(i)|v_i) = \prod_{j \in N(i)} P(v_j|v_i), \quad P(v_j|v_i) = \frac{\exp(u_j^T u_i)}{\sum_{k=1}^{|V|} \exp(u_k^T u_i)}$$

Random Walk: node2vec

Optimization problem [3]:

$$\max \sum_{v_i \in V} \log P(N(i)|v_i)$$

Conditional independence:

$$P(N(i)|v_i) = \prod_{j \in N(i)} P(v_j|v_i), \quad P(v_j|v_i) = \frac{\exp(u_j^T u_i)}{\sum_{k=1}^{|V|} \exp(u_k^T u_i)}$$

Simplified optimization problem:

$$\max \sum_{v_i \in V} \left[-\log Z_{v_i} + \sum_{j \in N(i)} u_j^T u_i \right], \quad Z_{v_i} = \sum_{v_j \in V} \exp(u_j^T u_i)$$

Random Walk: node2vec

Let c_i denote the i -th node in the random walk:

$$P(c_i = x | c_{i-1} = v) = \begin{cases} \frac{\pi_{vx}}{Z} & \text{if } (v, x) \in E \\ 0 & \text{otherwise} \end{cases}$$

Assume the walk just traversed edge (t, v) :

$$\pi_{vx} = \alpha_{pq}(t, x) \cdot w_{vx}, \quad \alpha_{pq}(t, x) = \begin{cases} \frac{1}{p} & \text{if } d_{tx} = 0 \\ 1 & \text{if } d_{tx} = 1 \\ \frac{1}{q} & \text{if } d_{tx} = 2 \end{cases}$$

Other: LINE

Joint probability between vertices [5]:

$$P(v_i, v_j) = \frac{1}{1 + \exp(-\mathbf{u}_i^T \mathbf{u}_j)}, \quad \hat{P}(v_i, v_j) = \frac{w_{ij}}{W}$$

Conditional distribution of the contexts:

$$P(v_j | v_i) = \frac{\exp(\mathbf{u}_j'^T \mathbf{u}_i)}{\sum_{k=1}^{|V|} \exp(\mathbf{u}_k'^T \mathbf{u}_i)}, \quad \hat{P}(v_j, v_i) = \frac{w_{ij}}{d_i}$$

Other: LINE

Joint probability between vertices [5]:

$$P(v_i, v_j) = \frac{1}{1 + \exp(-\mathbf{u}_i^T \mathbf{u}_j)}, \quad \hat{P}(v_i, v_j) = \frac{w_{ij}}{W}$$

Conditional distribution of the contexts:

$$P(v_j | v_i) = \frac{\exp(\mathbf{u}_j'^T \mathbf{u}_i)}{\sum_{k=1}^{|V|} \exp(\mathbf{u}_k'^T \mathbf{u}_i)}, \quad \hat{P}(v_j, v_i) = \frac{w_{ij}}{d_i}$$

First-order proximity loss:

$$O_1 = - \sum_{(i,j) \in E} w_{ij} \log p_1(v_i, v_j)$$

Second-order proximity loss:

$$O_2 = - \sum_{(i,j) \in E} w_{ij} \log p_2(v_j | v_i)$$

Other: LINE

Use negative sampling to optimize O_2 :

$$\log \sigma(u_j'^T u_i) + \sum_{i=1}^K E_{v_n \sim P_n(v)} [\log \sigma(-u_n'^T u_i)]$$

O_1 has a trivial minima, so we modify it to utilize negative sampling:

$$\log \sigma(u_j^T u_i) + \sum_{i=1}^K E_{v_n \sim P_n(v)} [\log \sigma(-u_n^T u_i)]$$

Other: LINE

Use negative sampling to optimize O_2 :

$$\log \sigma(u_j'^T u_i) + \sum_{i=1}^K E_{v_n \sim P_n(v)} [\log \sigma(-u_n'^T u_i)]$$

O_1 has a trivial minima, so we modify it to utilize negative sampling:

$$\log \sigma(u_j^T u_i) + \sum_{i=1}^K E_{v_n \sim P_n(v)} [\log \sigma(-u_n^T u_i)]$$

In case of low degree we add weight to second-order neighbors:

$$w_{ij} = \sum_{k \in N(i)} w_{ik} \frac{w_{kj}}{d_k}$$

Deep Learning: SDNE

Encoder-decoder model [6]:

$$y_i^{(k)} = \sigma(W^{(k)}y_i^{(k-1)} + b^{(k)}), \quad k = 1, \dots, K, \quad y_i^{(0)} = x_i$$

$$\hat{y}_i^{(k)} = \sigma(\hat{W}^{(k)}\hat{y}_i^{(k-1)} + \hat{b}^{(k)}), \quad k = 1, \dots, K, \quad \hat{y}_i^{(0)} = y_i^{(K)}, \quad \hat{y}_i^{(K)} = \hat{x}_i$$

Deep Learning: SDNE

Encoder-decoder model [6]:

$$y_i^{(k)} = \sigma(W^{(k)}y_i^{(k-1)} + b^{(k)}), \quad k = 1, \dots, K, \quad y_i^{(0)} = x_i$$

$$\hat{y}_i^{(k)} = \sigma(\hat{W}^{(k)}\hat{y}_i^{(k-1)} + \hat{b}^{(k)}), \quad k = 1, \dots, K, \quad \hat{y}_i^{(0)} = y_i^{(K)}, \quad \hat{y}_i^{(K)} = \hat{x}_i$$

Loss functions:

$$L_1 = \sum_{i,j=1}^n \alpha_{i,j} \|y_i^{(k)} - y_j^{(k)}\|_2^2$$

Deep Learning: SDNE

Encoder-decoder model [6]:

$$y_i^{(k)} = \sigma(W^{(k)}y_i^{(k-1)} + b^{(k)}), \quad k = 1, \dots, K, \quad y_i^{(0)} = x_i$$

$$\hat{y}_i^{(k)} = \sigma(\hat{W}^{(k)}\hat{y}_i^{(k-1)} + \hat{b}^{(k)}), \quad k = 1, \dots, K, \quad \hat{y}_i^{(0)} = y_i^{(K)}, \quad \hat{y}_i^{(K)} = \hat{x}_i$$

Loss functions:

$$L_1 = \sum_{i,j=1}^n a_{i,j} \|y_i^{(k)} - y_j^{(k)}\|_2^2$$

$$L_2 = \sum_{i=1}^n \|(\hat{x}_i - x_i) \odot b_i\|_2^2, \quad b_{i,j} = \begin{cases} 1 & \text{if } a_{i,j} = 0 \\ \beta > 1 \end{cases}$$

Deep Learning: SDNE

Encoder-decoder model [6]:

$$y_i^{(k)} = \sigma(W^{(k)}y_i^{(k-1)} + b^{(k)}), \quad k = 1, \dots, K, \quad y_i^{(0)} = x_i$$

$$\hat{y}_i^{(k)} = \sigma(\hat{W}^{(k)}\hat{y}_i^{(k-1)} + \hat{b}^{(k)}), \quad k = 1, \dots, K, \quad \hat{y}_i^{(0)} = y_i^{(K)}, \quad \hat{y}_i^{(K)} = \hat{x}_i$$

Loss functions:

$$L_1 = \sum_{i,j=1}^n \alpha_{i,j} \|y_i^{(k)} - y_j^{(k)}\|_2^2$$

$$L_2 = \sum_{i=1}^n \|(\hat{x}_i - x_i) \odot b_i\|_2^2, \quad b_{i,j} = \begin{cases} 1 & \text{if } \alpha_{i,j} = 0 \\ \beta > 1 \end{cases}$$

$$L = L_2 + \alpha L_1 + \nu L_{\text{reg}}, \quad L_{\text{reg}} = \frac{1}{2} \sum_{k=1}^K (\|W^{(k)}\|_F^2 + \|\hat{W}^{(k)}\|_F^2)$$

Datasets

- **Blogcatalog, Flickr and Youtube**

Social networks of online users. Each user is labelled by at least one category.

- **Arxiv GR-QC**

Paper collaboration network which covers papers in the field of General Relativity and Quantum Cosmology from arXiv.

- **20-Newsgroup**

Tf-idf vectors of each word are used to represent documents. The documents are connected based on their cosine similarity.

Datasets

Dataset	$ V $	$ E $
Blogcatalog	10312	667966
Flickr	80513	11799764
Youtube	1138499	5980886
Arxiv GR-QC	5242	28980
20-Newsgroup	1720	Full-connected

Experiment: reconstruction task

Arxiv GR-QC [6]

	SDNE	GraRep	LINE	DeepWalk
MAP	0.836	0.05	0.69	0.58

Blogcatalog [6]

	SDNE	GraRep	LINE	DeepWalk
MAP	0.63	0.42	0.58	0.28

Experiment: link prediction

Arxiv GR-QC [6]

Algorithm	P@2	P@10	P@100	P@200	P@300
SDNE	1	1	1	1	1
LINE	1	1	1	1	0.99
DeepWalk	1	0.8	0.6	0.555	0.443
GraRep	1	0.2	0.04	0.035	0.033

Summary

Summary

- Use matrix factorization when you can.

Summary

- Use matrix factorization when you can.
- Use random walk when matrix size is too large.

Summary

- Use matrix factorization when you can.
- Use random walk when matrix size is too large.
- Use deeper models on top of rough embeddings.

Summary

- Use matrix factorization when you can.
- Use random walk when matrix size is too large.
- Use deeper models on top of rough embeddings.

Follow us on github:

<https://github.com/src-d/role2vec>

References i

[1] S. Cao, W. Lu, and Q. Xu.

Grarep: Learning graph representations with global structural information.

In Proceedings of the 24th ACM International on Conference on Information and Knowledge Management, CIKM '15, pages 891–900, New York, NY, USA, 2015. ACM.

[2] P. Goyal and E. Ferrara.

Graph embedding techniques, applications, and performance: A survey.

arXiv preprint arXiv:1705.02801, 2017.

[3] A. Grover and J. Leskovec.

node2vec: Scalable feature learning for networks.

In Proceedings of the 22nd ACM SIGKDD international conference on Knowledge discovery and data mining, pages 855–864. ACM, 2016.

References ii

- [4] M. Ou, P. Cui, J. Pei, Z. Zhang, and W. Zhu.
Asymmetric transitivity preserving graph embedding.
In *KDD*, pages 1105–1114, 2016.
- [5] J. Tang, M. Qu, M. Wang, M. Zhang, J. Yan, and Q. Mei.
LINE: Large-scale Information Network Embedding.
ArXiv e-prints, Mar. 2015.
- [6] D. Wang, P. Cui, and W. Zhu.
Structural deep network embedding.
In *Proceedings of the 22nd ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 1225–1234. ACM, 2016.